

I. INTRODUCTION

0. Overview

0.0. Statistical learning

- (Making some hypotheses before building models)
- Preferring numerical data
- e.g. SVM and instance-based learning (memory/lazy learning)

0.1. Symbolic learning

- (Rule-based learning)
- Requiring inherent finite values
- Preferring categorical data (unordered)
- e.g. decision trees and rule induction learning -> They require discretized data or have their own discretization process

1. Discretization

1.0. Overview

- One of the most effective data pre-processing technique in Data Mining
- Translating quantitative data into qualitative data
- Procuring a non-overlapping division of a continuous domain
- Ensuring an association between each numerical value and a certain interval

-> Diminishing data from a large domain of numeric values to a subset of categorical values -> data reduction mechanism

- **No free lunch**

- Classical data reduction methods are not expected to scale well when managing huge data - both in number of features and instances -> Need distributed version
- Although many state-of-the-art DM algorithms have been implemented in **MLlib (a part of Spark)**, it is **not the case for discretization** algorithms yet.

1.1 Main Objective:

- Presenting a **distributed version** of the **entropy minimization discretizer** using **Apache Spark**, which is based on **Minimum Description Length Principle (MDLP)**
- proving that well-known discretization algorithms as MDLP can be parallelized in these frameworks, providing good discretization solutions for Big Data analytics
- Transforming the iterativity yielded by the original proposal in a single-step computation.

II. BACKGROUND AND PROPERTIES

1. Discretization process

- Assuming a data set S consisting of N examples, M attributes and c class labels, a discretization scheme D_A would exist on the continuous attribute $A \in M$, which partitions this attribute into k discrete and disjoint intervals:

$$\{[d_0, d_1], (d_1, d_2], \dots, (d_{k_A-1}, d_{k_A}]\},$$

where d_0 and d_{k_A} are, respectively, the minimum and maximal value, and $P_A = \{d_1, d_2, \dots, d_{k_A-1}\}$ represents the set of **cut points** of A in **ascending order**.

- A typical **discretization process** generally consists of **four steps**:

- (1) **Sorting** the continuous values of the feature to be discretized
- (2) **Evaluating** a cut point for splitting or adjacent intervals for merging
- (3) **Splitting or merging** intervals of continuous values according to some defined criterion
- (4) **Stopping** at some point

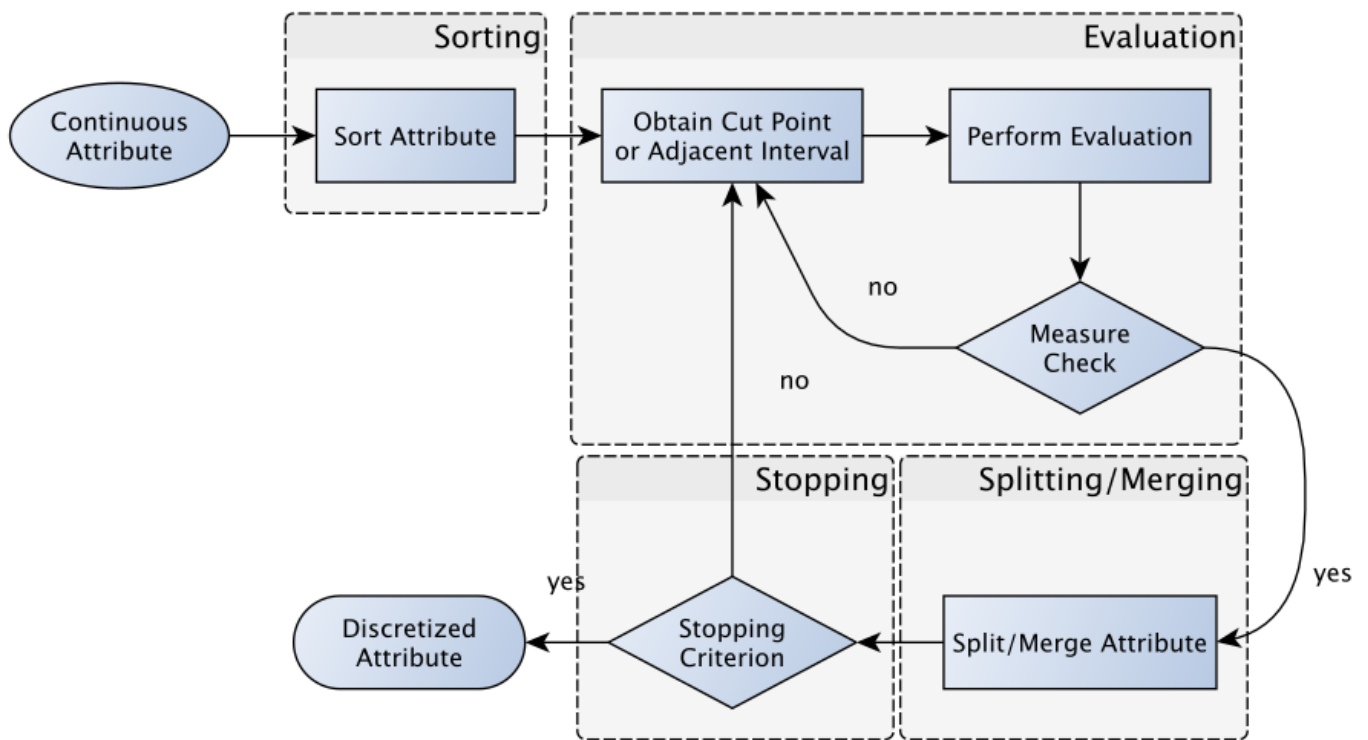


Figure 1: Discretization Process

- Sorting:** The continuous values for a feature are sorted in either **descending or ascending** order, with a time complexity of $O(N \log N)$
- Selection of a Cut Point:** After sorting, the **best cut point** or the **best pair of adjacent intervals** should be found in the attribute range in order to **split or merge** in a following required step. An **evaluation measure or function** is used to determine the **correlation, gain, improvement in performance** or **any other benefit** according to the **class label**.
- Splitting/Merging:** For **splitting**, the **possible cut points** are the **different real values** present in an attribute. For **merging**, the discretizer aims to **find the best adjacent intervals to merge** in **each iteration**.

- **Stopping Criteria:** When to stop the discretization process, **trade-off** between a **final lower number of intervals**, **good comprehension and consistency**.

2. Discretization Properties

2.1. Static vs. Dynamic:

- **Level of independence** between the **discretizer** and the **learning method**.
- A **static discretizer** is **run prior** to the **learning task** and is **autonomous** from the **learning algorithm**, as a **data pre-processing** algorithm.
- A **dynamic discretizer** responds when the learner requires so, **during the building of the model** -> **embedded in the learner** itself, producing **accurate and compact outcome** together with the associated learning algorithm.

2.2. Univariate vs. Multivariate

- **Univariate discretizers** only operate with a **single attribute simultaneously** -> **sort** the attributes **independently**, and then, the **derived discretization** disposal for **each attribute** keeps **unchanged** in the **next phases**.
- **Multivariate:** concurrently consider **all or various attributes**, may accomplish discretization **handling the complex interactions among several attributes** to decide also the attribute in which the **next cut point** will be **split or merged**.

2.3. Supervised vs. Unsupervised

- **Supervised** discretizers consider the **class label** whereas **unsupervised** ones do not.
- **Unsupervised discretization** can be applied on both supervised and unsupervised learning
- There is a **growing interest** in **unsupervised discretization** for **descriptive tasks**
- **Unsupervised** also open the door to **transfer the learning between tasks** since the discretization is not tailored to a specific problem.

2.4. Splitting vs. Merging

- **Splitting (Top down):** search for a **cut point** to **divide** the domain into **two intervals** among **all** the possible **boundary points**.
- **Merging (Bottom up):** begin with a **pre-defined partition** and search for a candidate **cut point** to **mix both adjacent intervals after removing it**.
- **Hybrid category:** as the way of **alternating splits with merges** during running time

2.5. Global vs. Local

- consider either **all available information** in the **attribute** or **only partial information**.
- **all the dynamic discretizers** and **some top-down** based methods are **local** (e.g. MDLP and ID3)

2.6. Direct vs. Incremental

- **Direct discretizer:** the **range associated to an interval** must be **divided into k intervals simultaneously**, requiring an **additional criterion** to **determine** the value of **k**. (e.g. **One-step discretization** methods)

- **Incremental methods:** begin with a **simple discretization** and pass through an **improvement process**, requiring an **additional criterion** to determine when it is the **best moment to stop**.

2.7. Evaluation Measure

- **Information:** Entropy, Gini index, Mutual Information
- **Statistical:** measurement of **dependency/correlation among attributes** ((Zeta, ChiMerge, Chi2), interdependency [40], probability and bayesian properties (MODL), contingency coefficient, etc.
- **Rough Sets:** evaluate the discretization schemes by **using rough set properties and measures**, such as **class separability, lower and upper approximations**, etc.
- **Wrapper:** rely on the **error** provided by a **classifier or a set of classifiers** that are used in each evaluation.
- **Binning:** there is **not an evaluation measure**. This refers to **discretize an attribute** with a **predefined number of bins** in a **simple way**. A bin assigns a **certain number of values per attribute** by using a **non sophisticated procedure**. (e.g. **EqualWidth** and **EqualFrequency** are **unsupervised** binning methods)

3. Minimum Description Length-based Discretizer (MDLP)

- Dynamic, univariate, supervised, splitting, local and incremental method
- Using the **Minimum Description Length Principle** to **control the partitioning process**.
- Introducing an **optimization** based on a **reduction of whole set of candidate points**, only formed by the **boundary points** in this set.

Denotation:

- $A(e)$: value for attribute A in the example e .
- b : A **boundary point** $\in Dom(A)$, which is the **midpoint** between $A(u)$ and $A(v)$, assuming that in the **sorted collection** of points in A , there exist two examples $u, v \in S$ with **different class labels**, such that $A(u) < b < A(v)$; and there does **not exist** other example $w \in S$ such that $A(u) < A(w) < A(v)$.
- B_A : set of **boundary points** for attribute A
- b_a : a boundary point to evaluate
- $S_1 \subset S$: a subset where $\forall a' \in S_1, A(a') \leq b_a$, and S_2 be equal to $S - S_1$.
- $E(S)$: class entropy in set S
- c, c_1 and c_2 : the number of class labels in S, S_1 and S_2 , respectively
- $N = |S|$: the number of training examples in set S

MDLP algorithm:

- **Recursively** evaluates **all boundary points**, computing the **class entropy of the partitions** derived as quality measure.
- The **objective is to minimize this measure** to obtain the **best cut decision**. The **class information entropy** yielded by a given **binary partitioning** can be expressed as:

$$EP(A, b_a, S) = \frac{|S_1|}{|S|} E(S_1) + \frac{|S_2|}{|S|} E(S_2)$$

- **Finally**, a **decision criterion** is defined in order to **control when to stop the partitioning process**. The use of **MDLP** as a **decision criterion** allows us to **decide whether or not to partition**. Thus a cut

point b_a will be applied iff:

$$G(A, b_a, S) > \frac{\log_2(N-1)}{N} + \frac{\Delta(A, b_a, S)}{N}$$

where $\Delta(A, b_a, S) = \log_2(3^c) - [cE(S) - c_1E(S_1) - c_2E(S_2)]$ and $G(A, b_a, S) = E(S) - EP(A, b_a, S)$

Important improvements:

- The number of cut points to derive in each iteration.
- A multi-interval extraction of points demonstrating that better classification models - both in error rate and simplicity.

III. TAXONOMY

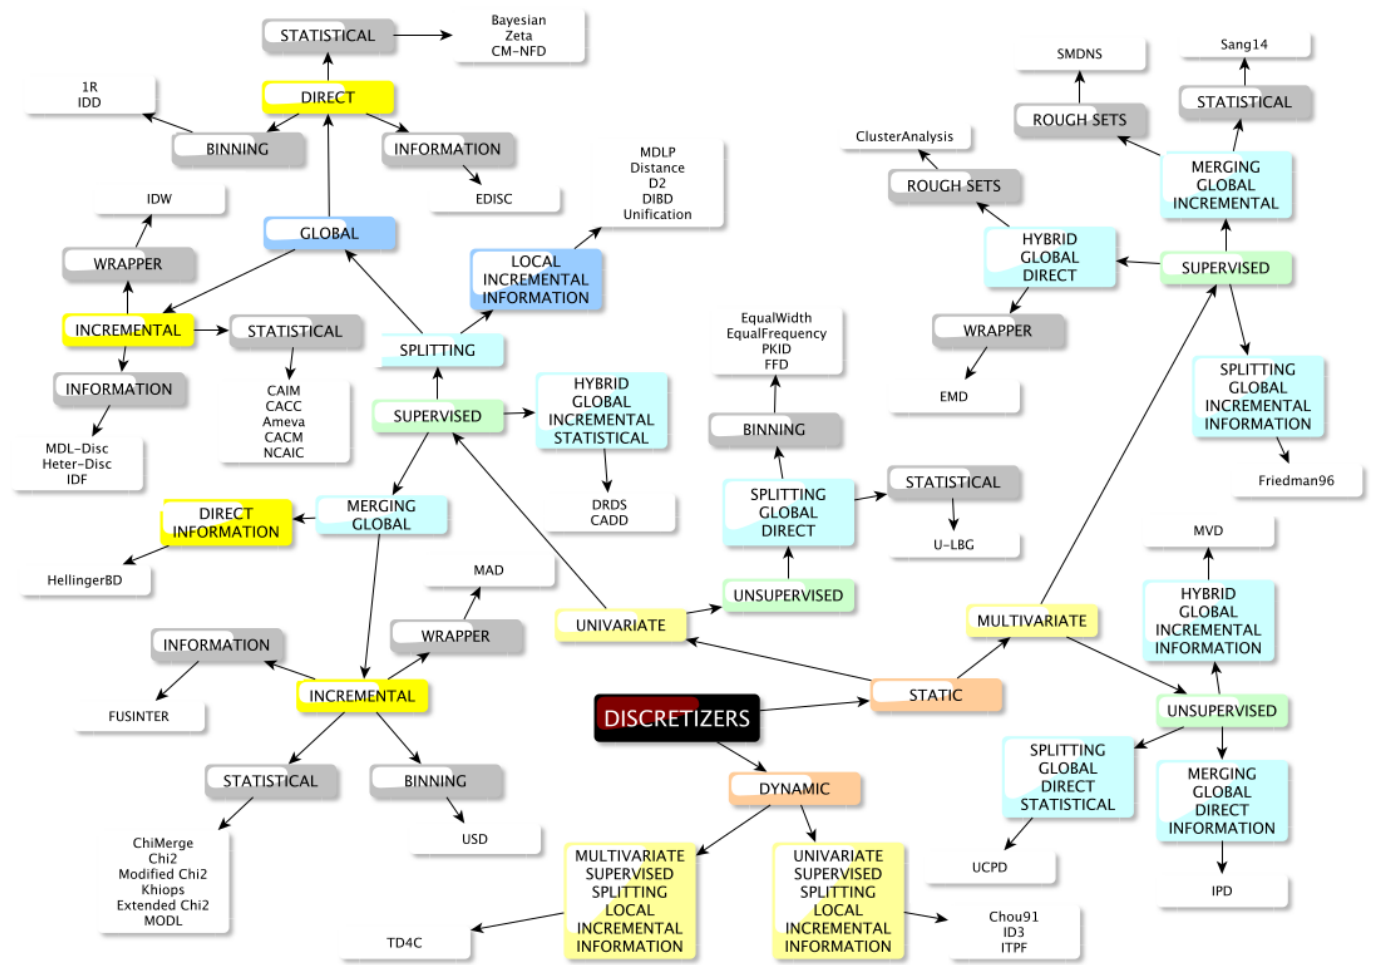


Figure 2: Discretization Taxonomy

- The **purpose of this taxonomy** is **three-fold**.
- **Firstly**, it **identifies the subset** of most representative **state-of-the-art discretizers**
- **Secondly**, it **characterizes the relationships** among techniques, the **extension of the families** and **possible gaps** to be filled in future developments.
- **No relevant methods** in the field of **Big Data**

IV. BIG DATA BACKGROUND

- The **3Vs terms** (extended with **2 additional terms**) that define it as **high volume, velocity and variety information** that require a **new large- scale processing**.
- **Volume:** the massive amount of data that is produced every day is still exponentially growing (from terabytes to exabytes)
- **Velocity:** data needs to be loaded, analyzed and stored as quickly as possible
- **Variety:** data come in many formats and representations
- **Veracity:** the quality of data to process
- **Value:** extracting value from data

MapReduce Model and Other Distributed Frameworks

- Aiming at **processing and generating large-scale datasets, automatically processed** in an **extremely distributed fashion** through **several machines**.
- Defining **two primitives** to work with **distributed data: Map and Reduce**.
- **Map: Splitting** data into **tuples (key/value pairs)** and **distributing** them across **clusters**
- **Reduce: Combining** those coincident **pairs** to form final output
- **Apache Hadoop -> Apache Spark** (based on **distributed data structures** called **Resilient Distributed Datasets (RDDs)**: immutable, partitioned collection of elements that can be operated on in parallel)

V. DISTRIBUTED MDLP DISCRETIZATION

- Distribute the complexity of this algorithm across the cluster.
- Two time-consuming operations: 1) Sorting of candidate points $O(|A|\log(|A|))$ complexity (assuming that all points in A are distinct) and 2) Evaluation of these points $O(|B_A|^2)$ -> both are bounded for a single attribute -> Algorithm to **sort and evaluate all points in a single step** -> Using some primitives from Spark's API, such as: **mapPartitions, sortByKey, flatMap and reduceByKey**.

1. Main discretization procedure

Algorithm 1: Main discretization procedure

- Calculates the **minimum-entropy cut points by feature** according to the **MDLP criterion**. It uses a parameter to **limit the maximum number of points** to yield.
- **Step 1: Generates tuples** with the value and the index for each feature as key and a class counter as value ($\langle A, A(s), v \rangle$).
- **Step 2: Reduces tuples** using a function that **aggregates** all **subsequent vectors** with the **same key**, obtaining the **class frequency** for **each distinct value** in the dataset.
- **Step 3: Sorts tuples by key** -> obtains the **list of distinct values** ordered by **feature index and feature value**.
- **Step 4:** Calculates the **first point by partition**
- **Step 5:** Generates the **boundary points**
- **Step 6:** Transforms **boundary points** into **tuples** with **feature index as key** ($\langle att, (point, q) \rangle$)
- **Step 7: Divides the tuples** in two groups (**big** and **small**) depending on **the number of candidate points by feature** exceeding a **threshold (mc)** or not and **re-formatted tuples** as: ($\langle point, q \rangle$).

- **Step 8: Evaluates and selects the most promising cut points grouped by feature (finally one best cut point per feature) according to the MDLP criterion (single-step version).**
- **Step 9: Joins both sets of cut points into final result**

Algorithm 1 Main discretization procedure

<p>Input: S Data set</p> <p>Input: M Feature indexes to discretize</p> <p>Input: mb Maximum number of cut points to select</p> <p>Input: mc Maximum number of candidates per partition</p> <p>Output: Cut points by feature</p> <pre> 1: $comb \leftarrow$ 2: map $s \in S$ 3: $v \leftarrow \text{zeros}(c)$ 4: $ci \leftarrow \text{class_index}(v)$ 5: $v(ci) \leftarrow 1$ 6: for all $A \in M$ do 7: $EMIT < (A, A(s)), v >$ 8: end for 9: end map </pre>	<pre> 10: $distinct \leftarrow \text{reduce}(comb, \text{sum_vectors})$ 11: $sorted \leftarrow \text{sort_by_key}(distinct)$ 12: $first \leftarrow \text{first_by_part}(sorted)$ 13: $bds \leftarrow \text{get_boundary}(sorted, first)$ 14: $bds \leftarrow$ 15: map $b \in bds$ 16: $< (att, point), q > \leftarrow b$ 17: $EMIT < (att, (point, q)) >$ 18: end map 19: $(small, big) \leftarrow \text{divide_attributes}(bds, mc)$ 20: $sth \leftarrow \text{select_thresholds}(small, mb, mc)$ 21: for all $a \in \text{keys}(big)$ do 22: $bth \leftarrow bth + \text{select_thresholds}(big(a), mb, mc)$ 23: end for 24: $\text{return}(\text{union}(bth, sth))$ </pre>
--	--

2. Boundary points selection

Algorithm 2: Function to generate the boundary points (get_boundary)

- Selects those points falling in the class borders, executes an independent function on each partition in order to parallelize the selection process as much as possible so that a subset of tuples is fetched in each thread.
- **Step 1:** For **each instance in a partition**, if the **feature index \neq index of the previous point** -> emit a **tuple** with the **last point as key** and the **accumulated class counter as value**. If not, checks whether the **current point** is a **boundary w.r.t the previous point** or not -> emits a **tuple** with the **midpoint as key** and the **accumulated counter as value**.
- **Step 2:** Repeat the scheme in step 1 for the **last point in the current partition** and the **first point in the next partition**.
- **Step 3:** Joins all tuples into a new mixed RDD of boundary points (bds)

Algorithm 2 Function to generate the boundary points (*get_boundary*)

Input: *points* An RDD of tuples $\langle (att, point), q \rangle$, where *att* represents the feature index, *point* the point to consider and *q* the class counter.

Input: *first* A vector with all first elements by partition

Output: An RDD of points.

```

1: boundaries  $\leftarrow$ 
2:   map partitions  $part \in points$ 
3:    $\langle (la, lp), lq \rangle \leftarrow next(part)$ 
4:    $accq \leftarrow lq$ 
5:   for all  $\langle (a, p), q \rangle \in part$  do
6:     if  $a \neq la$  then
7:       EMIT  $\langle (la, lp), accq \rangle$ 
8:        $accq \leftarrow ()$ 
9:     else if is_boundary( $q, lq$ ) then
10:      EMIT  $\langle (la, (p + lp)/2), accq \rangle$ 
11:       $accq \leftarrow ()$ 
12:   end if
13:    $\langle (la, lp), lq \rangle \leftarrow \langle (a, p), q \rangle$ 
14:    $accq \leftarrow accq + q$ 
15: end for
16:  $index \leftarrow get\_index(part)$ 
17: if  $index < npartitions(points)$ 
18:   then
19:      $\langle (a, p), q \rangle \leftarrow first(index + 1)$ 
20:     if  $a \neq la$  then
21:       EMIT  $\langle (la, lp), accq \rangle$ 
22:     else
23:       EMIT  $\langle (la, (p + lp)/2), accq \rangle$ 
24:     end if
25:   else
26:     EMIT  $\langle (la, lp), accq \rangle$ 
27:   end if
28: end map
29: return(boundaries)

```

3. MDLP evaluation

- Features in **each group** are **evaluated differently**
- **Small features** are evaluated in a **single step** where **each feature** corresponds with a **single partition**.
- **Big features** (*less frequent*) are evaluated **iteratively** since **each feature** corresponds with a **complete RDD with several partitions**.
- In **both cases**, the *select_thresholds* function is applied to **evaluate and select the most relevant cut points by feature**.
- **Small features** -> *arr_select_ths*; **big features** -> *rdd_select_ths*

3.1. Algorithm 3: Function to select the best cut points for a given feature (*select_thresholds*)

- **Evaluates and selects the most promising cut points grouped by feature** according to the **MDLP criterion (single-step version)**.
- **Step 1: Select the best cut point** (minimum entropy \leftrightarrow maximum entropy gain) satisfying **MDLP** (done by *arr_select_ths* and *rdd_select_ths*)
- **Step 2: Add** this point to **result** list and the **current subset** is **divided into two new partitions** using this cut point.
- **Repeats** until **no partition** to evaluate or **the number of selected points (mb)** is fulfilled

Algorithm 3 Function to select the best cut points for a given feature (*select_thresholds*)

Input: *cands* A RDD/array of tuples ($\langle point, q \rangle$), where *point* represents a candidate point to evaluate and *q* the class counter.

Input: *mb* Maximum number of intervals or bins to select

Input: *mc* Maximum number of candidates to eval in a partition

Output: An array of thresholds for a given feature

```

1: st  $\leftarrow$  enqueue(st, (candidates, ()))
2: result  $\leftarrow$  ()
3: while  $|st| > 0 \ \& \ |result| < mb$  do
4:   (set, lth)  $\leftarrow$  dequeue(st)
5:   if  $|set| > 0$  then
6:     if type(set) = 'array' then
7:       bd  $\leftarrow$  arr_select_ths(set, lth)
8:     else
9:       bd  $\leftarrow$  rdd_select_ths(set, lth, mc)
10:    end if
11:    if bd  $\neq$  () then
12:      result  $\leftarrow$  result + bd
13:      (left, right)  $\leftarrow$  divide(set, bd)
14:      st  $\leftarrow$  enqueue(st, (left, bd))
15:      st  $\leftarrow$  enqueue(st, (right, bd))
16:    end if
17:  end if
18: end while
19: return(sort(result))

```

3.2. Algorithm 4: Function to select the best cut point according to MDLP criterion (single-step version) (*arr_select_ths*)

- **Accumulates frequencies** and then **selects the minimum-entropy candidate** for **small attributes**
- **Step 1:** obtains the **total class counter vector** by **aggregating all candidate vectors**.
- **Step 2:** obtains the **accumulated counters** for the **two partitions** generated by **each point** (done by **aggregating the vectors** from the **most-left point** to the **current one**, and from the **current point** to the **right-most point**) -> ($\langle point, q, lq, rq \rangle$)
- **Step 3:** **Evaluates** the candidates using the *select_best* function

Algorithm 4 Function to select the best cut point according to MDLP criterion (single-step version) (*arr_select_ths*)

Input: *cands* An array of tuples ($\langle point, q \rangle$), where *point* represents a candidate point to evaluate and *q* the class counter.

Output: The minimum-entropy cut point

```

1: total  $\leftarrow$  sum_freqs(cands)
2: lacc  $\leftarrow$  ()
3: for  $\langle point, q \rangle \in cands$  do
4:   lacc  $\leftarrow$  lacc + q
5:   freqs  $\leftarrow$  freqs + (point, q, lacc, total - lacc)
6: end for
7: return(select_best(cands, freqs))

```

3.3. Algorithm 5: Function that selects the best cut points according to MDLP criterion (RDD version) (*rdd_select_ths*)

- Explains the **selection process**; in this case for **"big" features** (more than one partition).
- **Step 1: Re-distributes** this RDD into *npart*
- **Step 2:** Computes the **accumulated counter by partition**.
- **Step 3: Aggregates the results (by partition)** to obtain the **total accumulated frequency** for the whole subset

- **Step 4:** Computes the **accumulated frequencies by point from both sides** Firstly, the process **accumulates the counter from all previous partitions to the current one**. Then it computes the **accumulated values for each inner point** -> (*< point, q, lq, rq >*)
- **Step 6:** **Evaluates** the candidates using the *select_best* function

Algorithm 5 Function that selects the best cut points according to MDLP criterion (RDD version) (*rdd_select_ths*)

Input: *cands* An RDD of tuples (*< point, q >*), where *point* represents a candidate point to evaluate and *q* the class counter.

Input: *mc* Maximum number of candidates to eval in a partition

Output: The minimum-entropy cut point

```

1: npart ← round(|cands|/mc)
2: cands ← coalesce(cands, npart)
3: totalpart ←
4:   map partitions partition ∈ cands
5:     return(sum(partition))
6:   end map
7: total ← sum(totalpart)
8: freqs ←
9:   map partitions partition ∈ cands
10:    index ← get_index(partition)
11:    ltotal ← ()
12:    freqs ← ()
13:    for i = 0 until index do
14:      ltotal ← ltotal + totalpart(i)
15:    end for
16:    for all < point, q > ∈ partition do
17:      freqs ← freqs +
18:        (point, q, ltotal + q, total - ltotal)
19:    end for
20:    return(freqs)
21:  end map
22: return(select_best(cands, freqs))

```

3.4. Algorithm 6: Function that calculates class entropy values and selects the minimum entropy cut point (*select_best*)

- **Evaluates** the **discretization schemes** yielded by **each point** by **computing the entropy for each partition generated**, also taking into account the **MDLP criterion**.
- From the **set of accepted points**, the algorithm **selects** the **one with the minimum class information entropy**.

Algorithm 6 Function that calculates class entropy values and selects the minimum-entropy cut point (*select_best*)

Input: *freqs* An array/RDD of tuples ($\langle \text{point}, q, lq, rq \rangle$), where *point* represents a candidate point to evaluate, *leftq* the left accumulated frequency, *rightq* the right accumulated frequency and *q* the class frequency counter.

Input: *total* Class frequency counter for all the elements

Output: The minimum-entropy cut point

```

1:  $n \leftarrow \text{sum}(\text{total})$ 
2:  $\text{totalent} \leftarrow \text{ent}(\text{total}, n)$ 
3:  $k \leftarrow |\text{total}|$ 
4:  $\text{accp} \leftarrow$ 
5: map  $\langle \text{point}, q, lq, rq \rangle \in \text{freqs}$ 
6:    $k1 \leftarrow |lq|; k2 \leftarrow |rq|$ 
7:    $s1 \leftarrow \text{sum}(lq); s2 \leftarrow \text{sum}(rq);$ 
8:    $\text{ent1} \leftarrow \text{ent}(s1, k1); \text{ent2} \leftarrow \text{ent}(s2, k2)$ 
9:    $\text{partent} \leftarrow (s1 * \text{ent1} + s2 * \text{ent2}) / s$ 
10:   $\text{gain} \leftarrow \text{totalent} - \text{partent}$ 
11:   $\text{delta} \leftarrow \log_2(3^k - 2) - (k * h_s - k1 * \text{ent1} - k2 * \text{ent2})$ 
12:   $\text{accepted} \leftarrow \text{gain} > ((\log_2(s-1)) + \text{delta}) / n$ 
13:  if  $\text{accepted} = \text{true}$  then
14:     $\text{EMIT} \leftarrow \langle \text{partent}, \text{point} \rangle$ 
15:  end if
16: end map
17:  $\text{return}(\text{min}(\text{accp}))$ 

```

VI. EXPERIMENTAL FRAMEWORK AND ANALYSIS

Hardware

- A cluster composed of twenty computing nodes and one master node, each node has: 2 processorsx Intel Xeon CPU E5-2620, 6 cores per processor, 2.00 GHz, 15 MB cache, QDR InfiniBand Network (40 Gbps), 2 TB HDD, 64 GB RAM.

Software

- Hadoop 2.5.0-cdh5.3.1 from Cloudera's open-source Apache Hadoop distribution⁵, Apache Spark and MLlib 1.2.0, 480 cores (24 cores/node), 1040 RAM GB (52 GB/node).

Data set

Table 2 Summary description for classification datasets

Data Set	#Train Ex.	#Test Ex.	#Atts.	#Cl.
epsilon	400 000	100 000	2000	2
ECBDL14 (ROS)	65 003 913	2 897 917	631	2

- ECBDL14 is a **binary classification** problem where the **class distribution** is **highly imbalanced: 2% of positive instances**

Algorithms

- Handling imbalanced data distribution:** MapReduce version of the **Random OverSampling (ROS)** with **sampling_strategy = 1**

- **Classifier: Naive Bayes** (multinomial version in **MLlib** with $\lambda = 1$)

Evaluation Metrics

- **Classification accuracy** -> it is **not a proper metric** because of **imbalanced distribution** -> it should be **F1-measure** if 2 classes are **equally important** or **F2-measure** when **positive class is more costly**

Table 3 Classification accuracy values

Dataset	NB		NB-disc	
	<i>Train</i>	<i>Test</i>	<i>Train</i>	<i>Test</i>
<i>ECBDL14</i>	0.5260	0.6276	0.6659	0.7260
<i>epsilon</i>	0.6542	0.6550	0.7094	0.7065

- **Classification time** (NB vs. NB disc) and **discretization time** (sequential vs. distributed) (in seconds)

Table 4 Classification and discretization time values: with vs. w/o discretization -first two columns-, and sequential vs. distributed -last two columns- (in seconds)

Dataset	NB	NB-disc	Sequential	Distributed
<i>ECBDL14</i>	31.06	26.39	295 508	1 087
<i>epsilon</i>	5.72	4.99	5 764	476